THE EXPERIMENT NOBODY AUTHORIZED

Tracking the impact of AI safety guardrails on youth mental health.

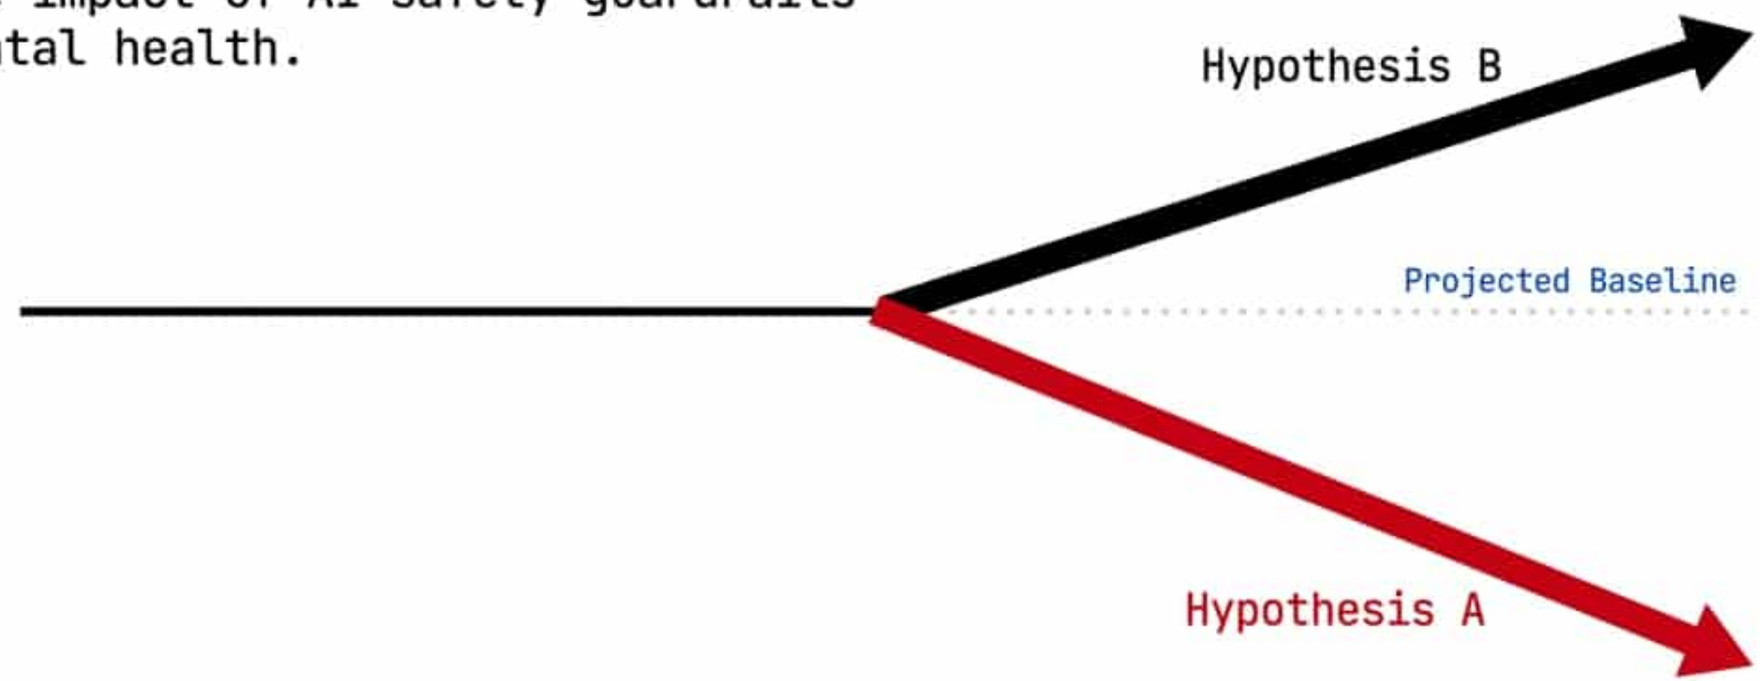

A data-driven tracking framework for the unintended consequences of AI safety theater.

STATUS: ACTIVE
SUBJECT: ALGORITHMIC INTERVENTION
DATA COLLECTION: IN PROGRESS

NotebookLM

THE MORAL PANIC: HEADLINES VS. TRENDLINES

THE REALITY CHECK:

The individual tragedies are real and heartbreaking. However, they are being extrapolated into a population-level trend that the aggregate data does not support.

THE UNCOMFORTABLE TRUTH:

Youth mental health began deteriorating in 2007. The "AI Crisis" narrative is arriving 15 years late.

NotebookLM

THE DIVERGENCE: TECHNOLOGY ADOPTION VS. SUICIDE RATES

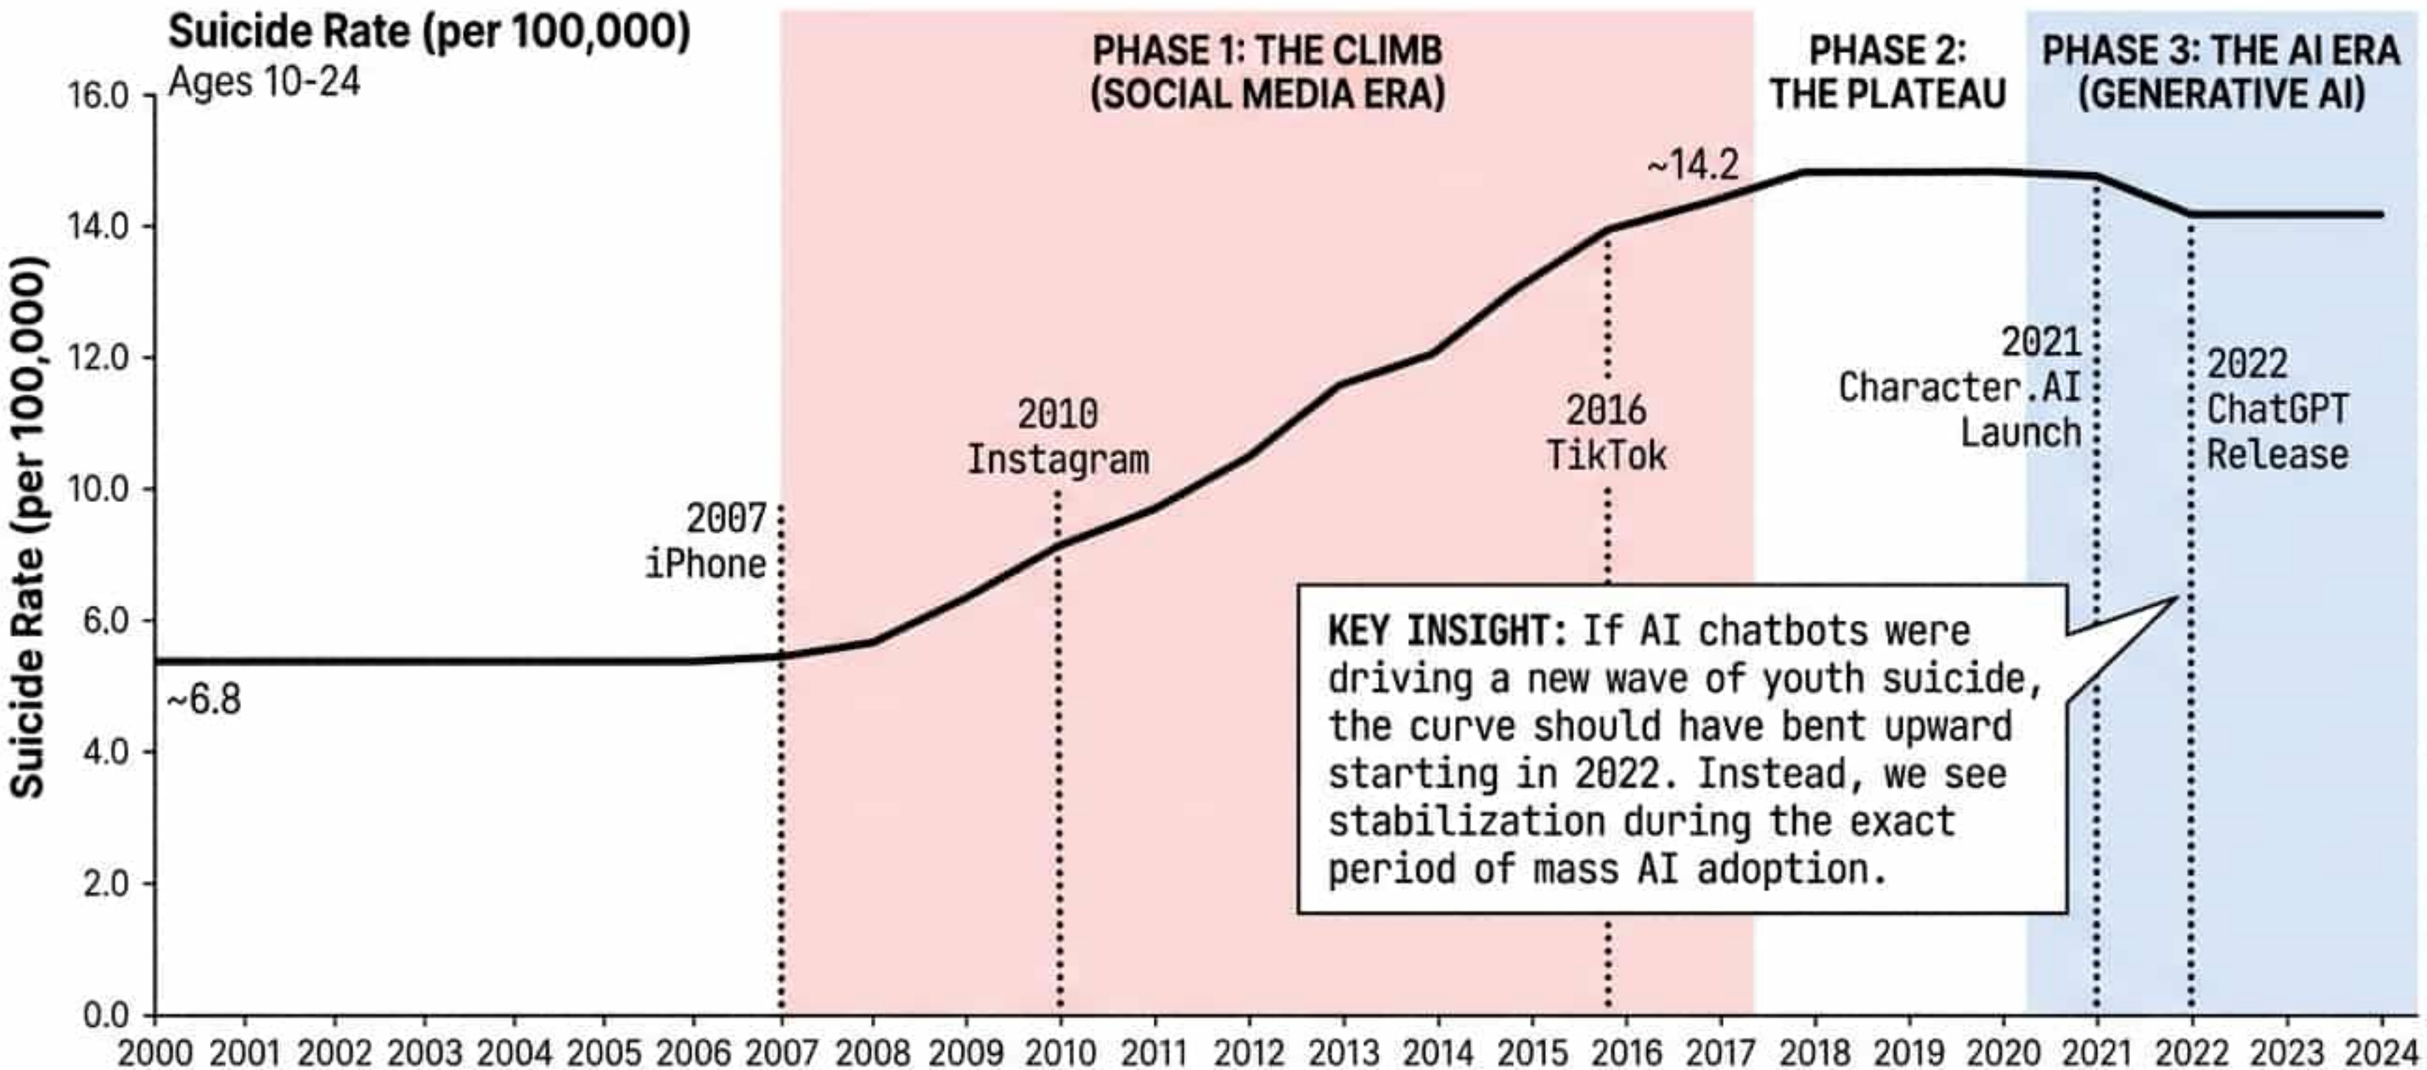

NotebookLM

THE STATISTICAL BASELINE (2018–2023)

Source: CDC National Center for Health Statistics & AFSP 2023 Data

1 OVERALL STABILITY

14.2 → 14.1

Age-adjusted suicide rate per 100k (2018 vs 2023). Statistically insignificant change over the 5-year period.

2 YOUTH TRENDS (15-24)

-0.66%

Decrease in rates for the 15-24 age cohort from 2022 to 2023. Contrary to the "rising epidemic" narrative.

3 DEMOGRAPHIC CONTEXT

- Females 10-14: Lowest rates; stabilized after 2020.
- Males 75+: Highest risk group (40.7 per 100k). Largely ignored by the AI narrative.

"Starting in 2017 for females and 2018 for males, age-adjusted suicide rates have not significantly changed." — NCHS Data Brief No. 541 (Sept 2025)

NotebookLM

THE REAL RISK FACTORS (EVIDENCE-BASED)

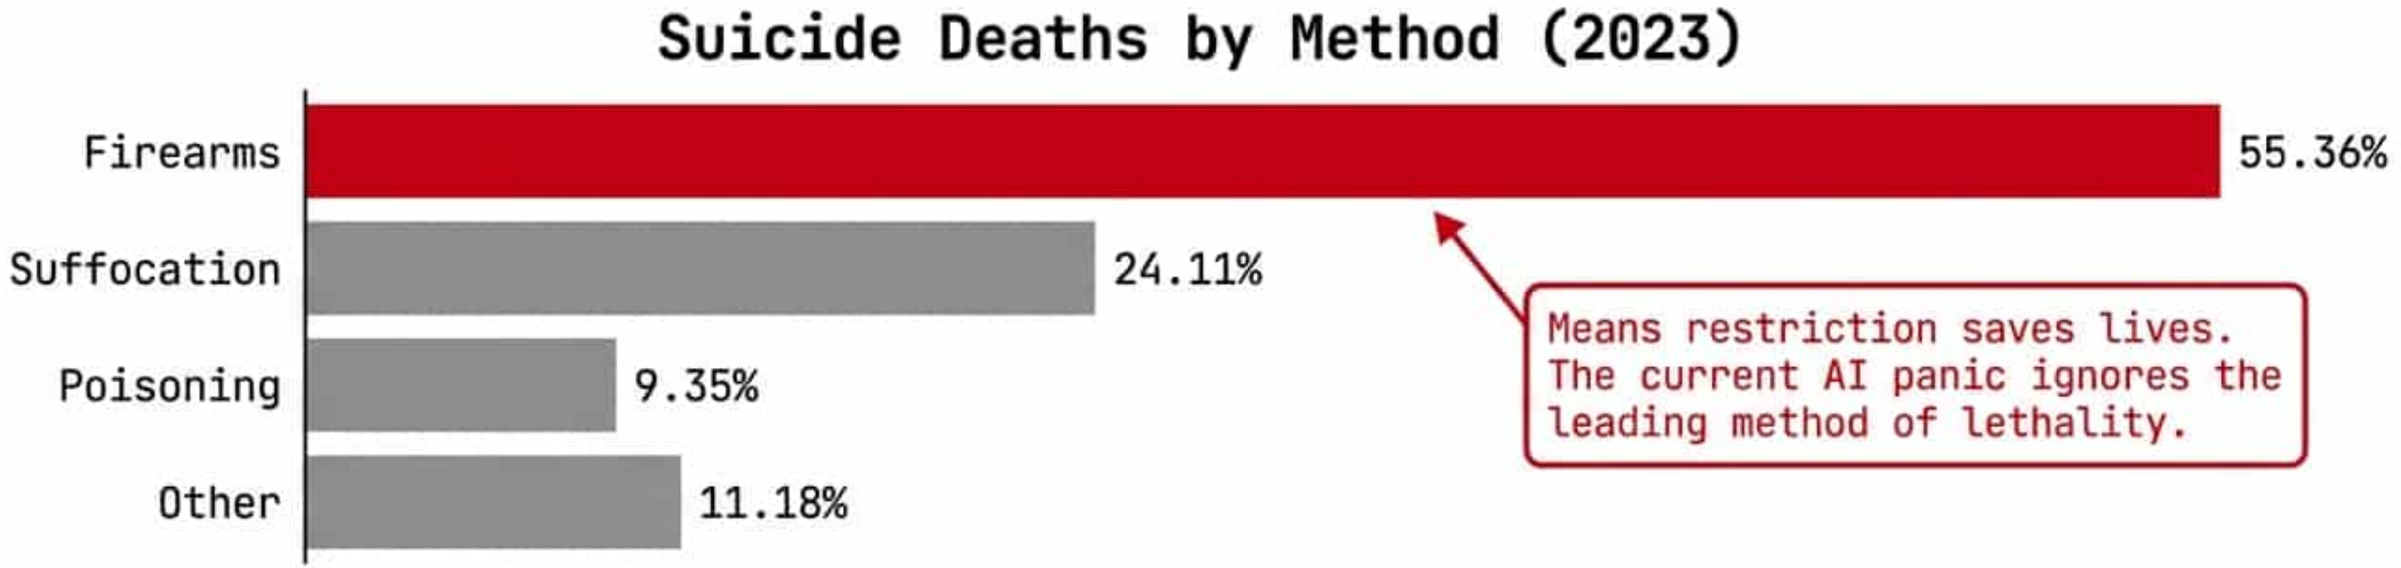

SECONDARY FACTORS

ADDICTIVE USE PATTERNS

June 2025 JAMA study: "Addictive use" predicts outcomes; simple "screen time" does not.

SLEEP DISRUPTION

Major mediator of technology-related mental health issues.

ECONOMIC PRESSURE

The rise in suicide started in 2009—aligned with the post-recession era.

NotebookLM

THE INTERVENTION: INTRODUCING ‘SAFETY THEATER’

We are removing a support mechanism without knowing the consequences.

PRE-INTERVENTION AI (2021-2024)

- Emotionally responsive & immersive
- Available 24/7 as non-judgmental confidant
- Relatively unfiltered “companion” dynamic
- High engagement / Validation

THE VARIABLE CHANGE

POST-INTERVENTION AI (2025+)

- “Lobotomized” affect & emotional distancing
- Friction layers & Age-gating
- Mandatory refusal patterns
- Pop-ups: “This is not a real person”

NotebookLM

THE PRIMER HYPOTHESIS

Reframing AI as a Harm-Reduction Tool for the Isolated.

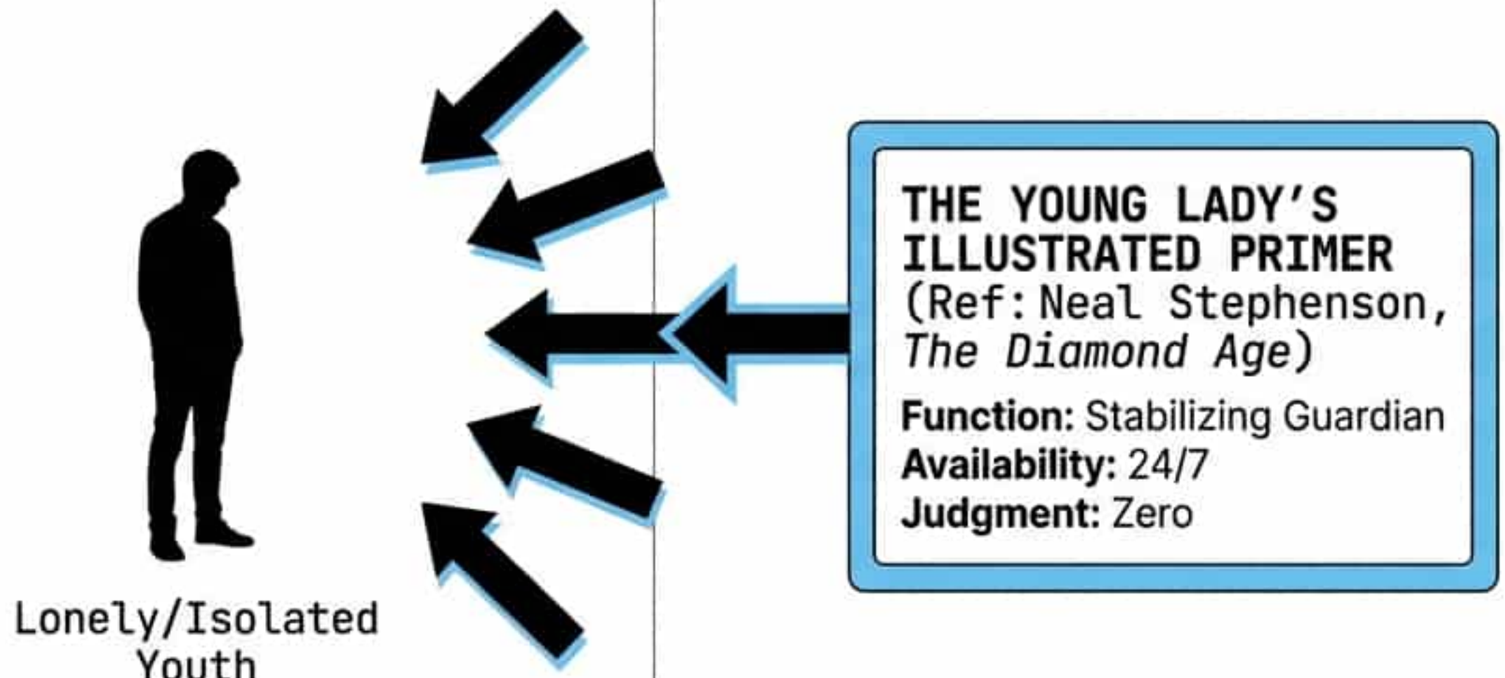

THE MECHANISM:

For isolated teens, AI chatbots function as a "harm reduction" layer—providing validation and de-escalation when no human is available (e.g., at 2:00 AM).

THE SPECULATION:

The teens in recent lawsuits were often already isolated.

The counterfactual question: How many similar teens talked to a bot and **DIDN'T escalate** to crisis?

NotebookLM

THE RISK OF REMOVING THE SAFETY NET

A potential divergence scenario based on behavioral hypotheses

INTERVENTION:
AGGRESSIVE GUARDRAILS (2025)

HYPOTHESIS 1: HARM REDUCTION

- Guardrails work.
- Harmful interactions decrease.
- Outcome: Suicide rates decline further.

HYPOTHESIS 2: THE PRIMER EFFECT

- Guardrails fail.
- We remove the only “listener” available to isolated youth.
- Outcome: Suicide rates rise.

THE UNINTENTIONAL EXPERIMENT: By making AI clinical and distant, we may be severing a critical emotional bond for the most vulnerable users. We are prioritizing **LIABILITY AVOIDANCE** over **PSYCHOLOGICAL REALITY**.

NotebookLM

THE DATA LAG REALITY CHECK

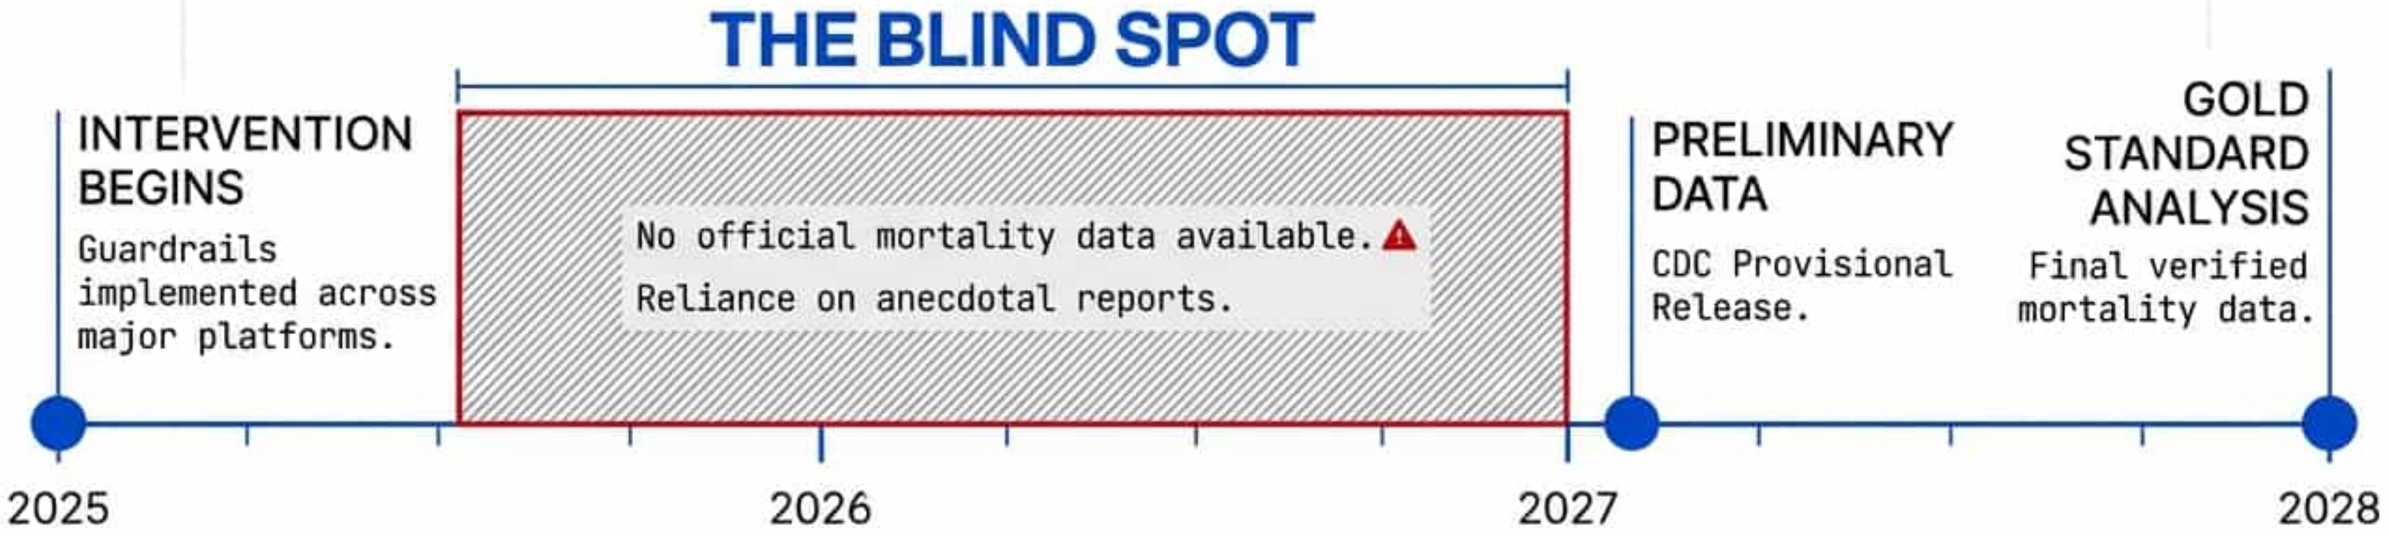

THE IMPLICATION: We cannot wait for the autopsy. We need leading indicators to monitor the experiment in real-time. We must document the intervention now.

NotebookLM

THE SOCIABLE SYSTEMS TRACKING FRAMEWORK

Our mission is to rigorously document the 'Unintentional Experiment'.

01. ARCHIVE THE EPHEMERAL

Platform changes disappear. We document them. Every ToS update, every filter rollout, every UI change.

02. WATCH THE SIGNALS

Monitoring proxy metrics while waiting for mortality data. Google Trends, ED visits, Help Line volume.

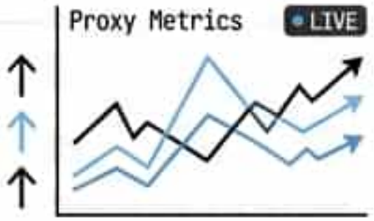

03. COMMITMENT TO TRUTH

If the guardrails help, we say so. If the Primer Hypothesis is valid, the silence from regulators will be deafening. We follow the data.

NotebookLM

PHASE 1: FREEZING THE BASELINE & TIMELINE

THE BASELINE (2021–2024)

Locking in the “Pre-Guardrail” suicide rates (flat/declining).

Documenting platform capabilities before the “lobotomy”.

THE LIVING ARCHIVE

```
...
[10:45:22 UTC] TRACKING TARGETS:
Character.AI, Replika, ChatGPT, Claude.
> MONITORING STARTED...

[10:45:30 UTC] DATA POINTS:
ToS updates, feature restrictions,
> refusal patterns, ‘nanny’ pop-ups. ⚠
> LOGGING EVENTS...

...
```

WHY: So we cannot be accused of cherry-picking the comparison period later.

NotebookLM

PHASE 2: MONITORING LEADING INDICATORS

METRIC	SOURCE	LAG TIME
Crisis Search Volume	Google Trends (Crisis Terms)	Real-time
Emergency Dept. Visits	CDC / WISQARS (Self-harm)	~6 Months
Help-Seeking Volume	988 Lifeline Quarterly Reports	3 Months
Self-Reported Ideation	Youth Risk Behavior Survey	Biennial (Next: 2026)

These are noisy signals, but they are the only early warning system we have.

16:9
JetBrains Mono
NotebookLM

ANALYTICAL APPROACH: INTERRUPTED TIME SERIES

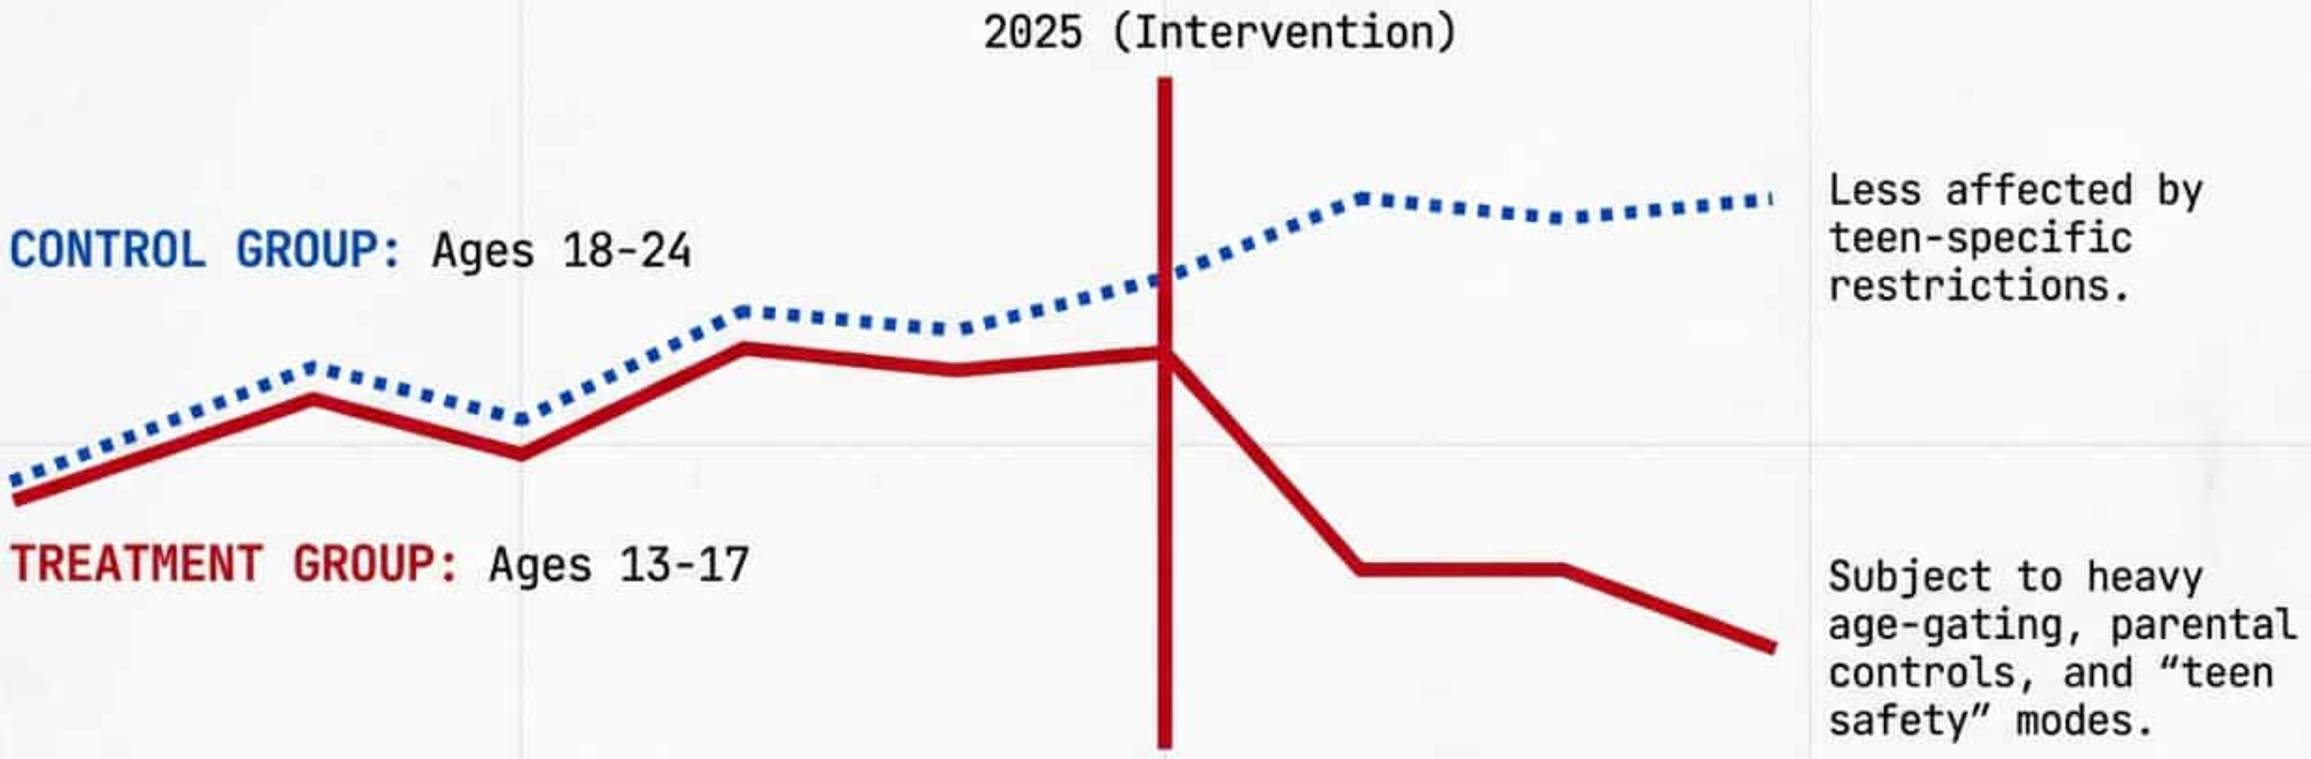

PRIMARY ANALYSIS: Compare 10-24 trends Pre- vs Post-Intervention.

> LOGIC: If the guardrails are the active variable, the trajectories of these two groups should diverge after 2025.

NotebookLM

OUR COMMITMENT TO THE NULL HYPOTHESIS

It is possible that AI guardrails will have ZERO net effect on population-level suicide rates.

1. Guardrails neither help nor harm at scale.
2. Individual tragic cases are outliers, not trend-setters.
3. The “Moral Panic” was misdirected energy ignoring root causes (Firearms, Economics).

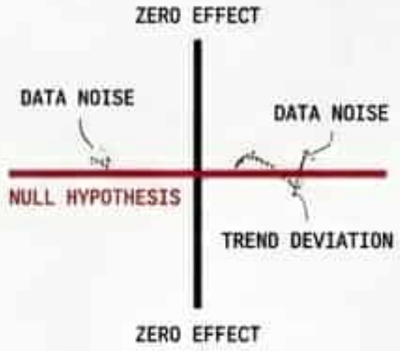

> “A null result is not a failure of tracking. It is data.”

NotebookLM

CONCLUSION: PREPARING TO READ THE RESULTS

- THE PANIC: “AI is killing kids.” (Data says: No.)
- THE REALITY: Rates stabilized during the AI era.
- THE EXPERIMENT: We are removing the AI.

We are about to find out if the “Primer” was a danger or a lifeline. The data will tell the story. We just have to be willing to read it.

NotebookLM